

Genetic Fuzzing

Biologically inspired artificial intelligence – a.y. 2025/26

Luca Sartore, Jacopo Sbabo, Alessio Zeni

Abstract—Assessing code robustness is a fundamental aspect of many software applications, as it enables more resilient implementations while also enhancing security by identifying unintended or spurious behaviors.

Robustness evaluation requires the definition of a structured testing suite, guided by one or more quantitative metrics, to systematically observe how a given codebase behaves under a variety of conditions.

A wide range of testing techniques exists, spanning from manual to fully automated approaches; these techniques differ in the aspects of the code they target (e.g., branch coverage, line coverage), in how inputs and outputs are generated or tracked (e.g., deterministic mappings versus random sampling), and in the feedback metrics they employ.

In this work, we focus on fuzzing, adopting Python as the reference programming language. Fuzzing is an automated testing technique that supplies random, malformed or unexpected inputs to a program with the goal of exposing erroneous or undefined behaviors.

The evolutionary component of our approach is realized through the use of novelty search, which drives the generation of inputs toward maximizing code exploration rather than optimizing for a specific failure condition.

To quantify the effectiveness of the proposed method, we rely on code coverage as the primary evaluation metric, leveraging the `coverage.py` library to measure both line and branch coverage.

Our results demonstrate that, although the proposed approach may perform comparably to simpler and faster techniques in worst-case scenarios, it is particularly effective when applied to deep and highly structured codebases. In such settings, it achieves superior coverage with fewer function calls compared to standard fuzzing strategies.

I. INTRODUCTION

CODE robustness is a cornerstone of programming; in many applications, especially security-dependent ones, reliability is fundamental to avoiding attacks and the exploitation of bugs. Standard applications also benefit from well-written and tested code, creating a smoother user experience and simplifying codebase maintenance.

For this reason, we must create test suites for our code, which generally aim for three different goals depending on requirements:

- **Bug-hunting:** looking for unwanted or anomalous behaviors
- **Requirements satisfaction:** ensuring the code satisfies the original constraints
- **Code coverage:** analyzing results from a wide variety of input combinations

Many testing techniques are available, divided into various categories depending on the level (single component, multiple components, complete software), approach (static, dynamic, passive), and the tester's point of view (black, white, gray

box). It is also important to define metrics to evaluate the effectiveness of a test suite, though these are inherently application-dependent.

Generally, there is no free lunch; each application needs a proper evaluation strategy that provides a supervision signal and works in tandem with a tailored testing technique.

A. Main focus

In our case, we focused on *fuzzing*, an automated testing technique that involves the use of random, invalid, and unexpected data, typically employed when testing programs with structured inputs. During this process, the program is monitored for crashes, failing assertions, and memory leaks. Inputs that are at the limit of validity are particularly valuable, as they effectively test corner cases and help uncover bugs.

B. The genetic twist

To generate inputs for the code under test, various techniques may be employed, though standard approaches typically rely on random input generators.

However, random generators possess a fundamental flaw: as the input space expands, the number of possible combinations becomes intractable, making it difficult for the generator to produce valid or compatible inputs for effective testing.

To address this, we leverage a genetic algorithm—specifically *novelty search*. Instead of maximizing a traditional fitness function, this technique prioritizes novelty, which measures how “new” a specific individual is compared to the previously explored population.

Novelty search is an excellent candidate for software testing; unlike objective-based approaches that may converge prematurely toward local optima and limit coverage, novelty search promotes exhaustive exploration. This allows the algorithm to discover new execution branches and expose corner cases while avoiding the pitfalls of premature convergence.

C. Evaluation

The evaluation of performance was conducted using `coverage.py`, a Python library that provides specialized tools for measuring coverage during the testing process. Consequently, this project focused primarily on code coverage as the central metric.

In particular, we selected two types of coverage that specifically benefit from the exploratory nature of novelty search:

- **Line coverage:** measures the proportion of executable code lines actually traversed during testing.
- **Branch coverage:** evaluates the decision points within the code (e.g., if-statements), ensuring that both true and false paths are executed.

II. RELATED WORK

The research landscape surrounding the intersection of genetic algorithms and fuzzing remains relatively sparse; while studies have emerged over the last few years, the topic has yet to see widespread adoption. In particular, our approach of combining novelty search with fuzzing appears to be one of the inaugural attempts in this specific direction.

Nevertheless, we have selected several key works to highlight diverse methodologies for integrating machine learning with fuzzing:

- **Learn & Fuzz (Godefroid et al., 2017):** represents one of the foundational efforts in leveraging machine learning models to assist the fuzzing process.
- **V-Fuzz (Li et al., 2019):** utilizes an evolutionary algorithm for input generation, specifically targeting areas of the code predicted to be vulnerable. This methodology is the most closely aligned with our own objectives.
- **DARWIN (Jauernig et al., 2020):** a contemporary approach that employs an evolution strategy to dynamically adapt mutation probability distributions during fuzzing.

III. OVERVIEW

The general architecture of our system can be seen in figure 1. The idea is to start from a function we want to test, and use the type annotations to understand the input that the function takes. With this input we can use a “type adapter” (an interface that implements crossover and mutation operations on top of existing Python types) to generate a valid input to the function.

We can then execute the input, and measure the coverage. This information can then be used by a “strategy” to mutate the input, and try to improve the coverage. The main strategy we used is Novelty Search, but we also tested others as baselines.

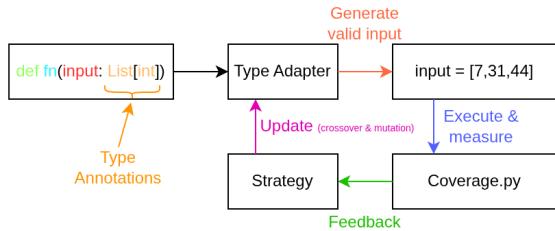


Fig. 1. Architecture overview

IV. IMPLEMENTATION

A. Tool Used

Our project is based on two main important tools:

1) *Coverage.py*: This is the primary python framework used for measuring code coverage. Code coverage is essential as it functions as our “objective function”.

2) *Type Annotations*: The “Type Annotations” feature of python allows a function to be marked with some metadata that indicate what kind of input the function expect. This information is used to select what kind of type adapter (see IV-B) should be used.

B. Type Adapters

Type adapters are an abstraction that implements genetic operators (namely crossover and mutation) on top of existing Python types (such as `int`, `bool`, `list` etc.). We implemented type adapters on top of the basic python types, and we defined an API to “inject” new adapters in the system (so that users of our library can use this feature to fuzz-test with custom types).

Our implementation of the type adapters is recursive, meaning that when we define the adapter for `list` we automatically get the adapters for all the specialized types such as `list[int]`, `extitlist[bool]`, `list[list[float]]` etc.

C. Dataset

We tried to use some existing datasets, but unfortunately none of the ones we found were adequate for our project. In particular we needed a dataset with type annotations, and we couldn’t find one. So we generated one using Google Gemini APIs. The dataset has approximately 100 functions and is fully open source.

D. Test Cases

For a complete overview of novelty search’s performance we created two other techniques to compare against: one that doesn’t use any particular logic (Random) and another that is a traditional Genetic Algorithm implementation (Input Bag).

1) *Random*: This strategy simply generate random inputs and throw them at the functions in order to measure the coverage.

2) *Input Bag*: Input Bag is our implementation of a traditional genetic algorithm.

Applying genetic algorithms to fuzzing has a fundamental problem: we need to evolve a set of inputs (not just one input) if we want to cover the entire function. To do so we had two options:

- **Bundling up multiple inputs in one individual** This idea is simple: we define a constant N that defines how many function inputs one individual is made of, and then we evolve the best N inputs that, when evaluated together, maximize coverage.
- **Using fitness sharing to encourage diversity** We can make each individual represent one single function input, and then use fitness sharing (or other similar techniques) to encourage individual diversity. In this way the entire population will become our final output (not only the fittest individual).

Both techniques were valid; we decided to use the first one for our work. We tested various selection/replacement strategies, and we found that, by far, the most important aspect was to use elitism (which makes sense given the highly non-linear nature of the problem). We also found that adding steady-state replacement to introduce novel genetic material helped a bit.

3) *Novelty Search*: Our implementation of novelty search is pretty standard (it keeps a list of all the “novel” individuals, and then uses them to generate newer ones). What is not straightforward is the way we define “Novelty.” The general

idea is to generate an hashable summary of the parts of the code hit by a certain input, so that we can understand if one specific input is “novel” with a simple hash-map (making this an $O(1)$ time operation). For example, in the case of line coverage the “summary” is a set containing all the lines that are executed. An input is considered “novel” if and only if the set of lines executed is different from all the others. This means that we could end up selecting individuals that cover less code than the ones we already have, assuming that they are doing something different. This behaviour is desired, as there are cases where reducing coverage can lead to exploring new paths, for example:

```
A = True
# not executing this if can increase
# coverage if C2 == True
if C1:
    A = False
if A and C2:
    to_something()
```

V. RESULTS

A. Setup

In our tests we evaluated the performance of the three strategies using the entire dataset. We ran two different test scenarios:

- **Function call parity:** in this scenario each strategy gets a limited amount of calls to the tested function, and we see which strategy works better with this constraint.
- **Execution time parity:** in this scenario we give each strategy a fixed execution time, and we allow them to work until the clock runs out.

In the table below (V-A) we can see that when using the “fixed execution time” setting, Novelty Search calls the test function about 400 times fewer than the other two strategies, which puts it at a significant disadvantage.

The reason for this behavior is that coverage.py has a significant overhead associated with starting a test. With *Random* and *Input Bag* we are able to amortize this cost by bundling multiple function calls into a single run; however, it is not possible to do the same using *Novelty Search*.

Given that this is more a limitation associated with coverage.py than something caused by *Novelty Search*, we decided to consider “Function call parity” for our final results, because otherwise *Novelty Search* would be unfairly disadvantaged.

Strategy	Total number of function calls	
	Fixed num of func calls	Fixed execution time
Random Search	10,000	~4,000,000
Input Bag	10,000	~4,000,000
Novelty Search	10,000	~10,000

B. General Trend

In the plot 2, we can see that Random Search and Input Bag perform similarly, with Random giving consistently more coverage than Input Bag. This is likely caused by more aggressive exploration in the Random Search approach, as well

as the fact that Input Bag is constrained in the number of inputs it can use.

Novelty Search outperforms both strategies in the end. We should note that the reason why Novelty Search lags behind in the beginning is that the first iteration of the other two strategies has already executed many more function calls due to the inputs being bundled.

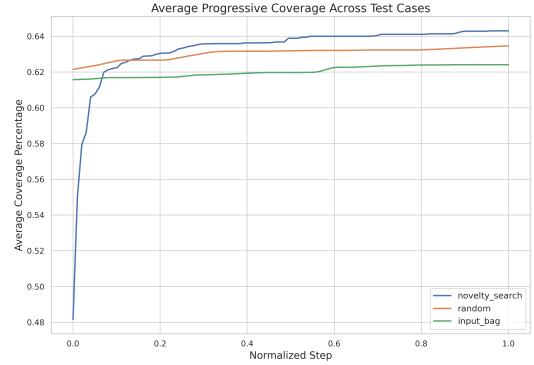


Fig. 2. Coverage progression over time

We had approximately 400 similar charts (for each individual function, and for both function call parity and execution time parity) that we can't include in the report for obvious reasons, but you can find them on [github](#) if you are interested.

C. A theoretical testcase: count_bool vs is_odd_for_dummies

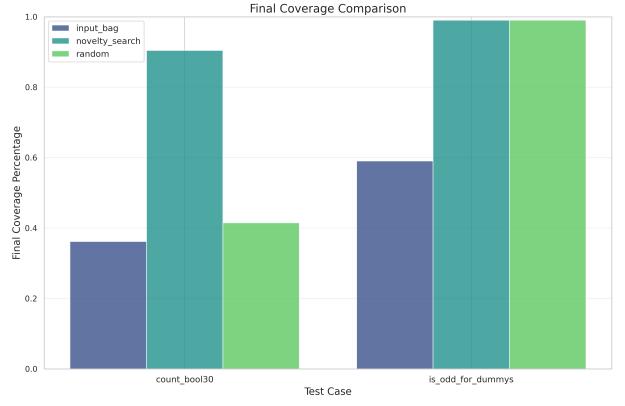


Fig. 3. count_bool vs is_odd_for_dummies

While testing the applications, we developed two pathological functions to test the theoretical strengths of each implementation.

- **count_bool:** this function takes a large number of boolean inputs and counts how many true values are present until the first false. We implemented it as a deep succession of if statements (code [here](#)).
- **is_odd_for_dummies:** this function takes an integer as input and has a long succession of if statements; each handles a particular integer, returning True if it is odd and False otherwise (code [here](#)).

The first function is very deep: to score a high output, all previous conditions must evaluate to True. Crafting an input that goes deep into the function is easier when done progressively (using evolution) instead of in one shot (using random input). We can see the results in plot 3. In the second case the function is very shallow, because to reach a particular branch the algorithm must generate an exact value, without any previous information to build it. This favors exploration over exploitation, reducing Novelty Search's advantage (as shown in plot 3).

D. The implications of the No Free Lunch Theorem

In our setup, we decided to test general arbitrary functions, since potentially each computable problem can be encoded in a function to test. Because of this, it is a fitting application for the No Free Lunch Theorem, which states that each optimization algorithm will perform on average the same over the space of all possible problems. This observation aligns with the discussion in section V-C.

E. Limitations of Coverage.py

We would like to note that the coverage.py framework had some limitations that impacted our results. The first is the overhead associated with starting a test, which we discussed in section V-A. However, we addressed that by switching to “function call parity” as the evaluation method.

The second limitation is perhaps harder to address. Coverage.py only measures line coverage and branch coverage, but there are other ways to measure coverage that could work better for our use case. For example, [Jest](#) (JavaScript’s testing framework) also measures *statement coverage*, which is more detailed and can provide a stronger supervisory signal. For example, take the following code:

```
if C1 and C2:
    do_something();
```

In the example we want to make C1 and C2 both true at the same time to cover the line inside the if statement. Using line coverage, a mutation that makes C1 true would go unnoticed, as it would not affect coverage in the immediate term. If we were instead using statement coverage, the mutation that makes C1 true would immediately result in higher coverage, and therefore a higher likelihood that the individual is selected for future generations and may then make C2 true in a subsequent mutation.

Another useful metric that coverage.py does not collect is the number of times that a certain line or branch is hit. Instead, coverage.py only collects information on whether a line is executed or not, ignoring any sort of counter. It is easy to see how this kind of information could be useful when testing code such as the following:

```
counter = 0
while condition():
    # a line counter here would help
    counter += 1
if counter > 3:
    do_something()
```

VI. CONCLUSIONS AND FUTURE WORK

In conclusion, using Novelty Search for fuzzing looks promising.

There are definitely some scenarios where Novelty Search outperforms our Random and Input Bag baselines by significant margins. However, there are also scenarios where Novelty Search is outperformed by even simple Random Search.

Two main challenges need to be addressed before the framework we developed could become useful in a production environment, and both are associated with coverage.py:

- **High overhead:** As we explained in section V-A, due to overheads associated with the initialization of coverage.py we had to use “function call parity” for all of our tests. However, in a production environment “execution time parity” is much more aligned with how the costs of various CI/CD pipelines are calculated, and therefore should be the main metric considered when evaluating performance.
- **Weak supervisory signal:** As we alluded to in section V-E, coverage.py only has limited ways to measure coverage (namely line and branch coverage). More complex (and potentially more effective) methods exist that could provide a stronger supervisory signal and therefore better guide our algorithm.

In future work we would like to address those limitations either by re-implementing parts of coverage.py to address its issues or, more realistically, by switching to a programming language that already has a better testing framework implemented.

It is also clear to us that, given the high generality of the concept of “a function” and given the implications of the No Free Lunch theorem, it will be impossible to have a single algorithm that can efficiently fuzz-test any kind of function. To address this limitation it would be interesting to test a “Mixture of Strategies” approach, where different strategies are executed on the same function, and over time only those that achieve the best results are kept while the others are discarded.

VII. INDIVIDUAL MEMBERS CONTRIBUTION

For the individual contributions, we divided the work as follows:

- Alessio and Jacopo read the paper on MAP-Elites and found it was not applicable to this work.
- Jacopo researched the literature and identified the novelty search approach, together with other possible approaches.
- Alessio implemented the code that measured the coverage and integrated our project with coverage.py.
- Luca was responsible for the generation of the graphs and the analytics.
- Luca and Alessio worked on the test runner that executed multiple tests using multithreading.
- All together, we worked on the genetic algorithm implementation: some implemented the outlines or the different adapter types, while others tested different values and approaches.